

Introduction to Pandas* (in Tulip):

*Panel Data Structures

Jocelyn Forest, Benjamin Renoust, Guy Mélançon,
LaBRI, Université Bordeaux I and CNRS UMR 5800

What is pandas?

- Recent API based on Numpy
- Devised by Wes McKinney
- Fast and intuitive data structures
- Easy to work with messy and irregularly indexed data
- Optimized for performance, with critical code paths compiled to C
- Adopts concepts of R language

Main focus

- The two basics structures of pandas
 - Series 1d array
 - DataFrame 2d array
 - Panel nd array ($n > 2$)
- Filtering, selecting data
- Aggregating, transforming data
- Joining, concatenating, merging data
- Descriptive basics statistics

Installing pandas

- Version python 2.6 or 2.7
- Dependencies:
 - NumPy 1.6.1 or higher
- Optional dependencies:
 - Matplotlib to plot
 - SciPy for statistical functions

Exercise

- ```
> sudo apt-get install python-pandas
```
- ```
> git clone git://github.com/pydata/pandas.git
```

```
> cd pandas
```

```
> python setup.py install
```
- Header:

```
import pandas as pd
```

Series

	index		value
0	C	▶	3
1	B	▶	7
2	A	▶	4
3	D	▶	4
4	D	▶	0.3

- Subclass from `numpy.ndarray`
- Any type of data (numeric, string, boolean...)
- Index need not to be ordered
- Duplicated index are possible

Some vocabulary:

- `Series.index`: list of indices
- `Series.values`: list of values

DataFrame

columns

index

	id	country	isOver	amount
a	P255	Afg	True	300000
b	P31256	Fr	False	22354
c	P2245	Cor	False	12478
d	415	Som	False	Nan
e	P332	Esp	True	4789123

- ndarray-like
- 2D data structure (for n D data structures see Panel)
- Dictionary of series
- Row and column index
- Size mutable: insert or delete columns

DataFrame

- Some vocabulary
 - `DataFrame.index`: list of DataFrame indices
 - `DataFrame.values`: 2D array of all values contained in the DataFrame
 - `DataFrame.columns`: list of columns labels
 - `axis`: indicates the axis index for rows (`axis = 0`), columns (`axis = 1`),
or even n th axis in panels

Panel

- Container for three or more dimensional data
- Dictionary of DataFrame objects
- Less used than Series and DataFrame objects
- DataFrame methods not all available for Panel objects
- **Unnecessary** in a lot of cases :
 - *Hierarchical indexing*

Construction of Series and DataFrame

Exercise

- Directly editing

```
s = pd.Series([3,7,4,4,0.3] ,  
              [index = ['a','b','c','d','e']])  
  
df = pd.DataFrame(np.arange(9).reshape(3,3),  
                  [index = ['b','a','c'],  
                   columns=['Paris','Berlin','Madrid']])
```

- From a python dict

```
data = {'Paris': [0,3,6,999999999],'Berlin': [1,4,7], 'Madrid': [2,5,8]}  
  
df = pd.DataFrame(data,  
                  [index = ['b','a','c','d'],  
                   Columns = ['Paris', 'Berlin', 'Madrid'] ])
```

Warning: index array size >= max element array size

- Several methods in the API to import from databases

```
df = pd.read_csv(path/fichier.csv,  
                 [index_col = [...]])  
  
df = pd.read_table(path/fichier.txt,  
                  [sep = ','])
```

Selection of data

- Selection on series

In: <code>s</code>		In: <code>s['b']</code>		In: <code>s['a':'c']</code>		In: <code>s['d']</code>		In: <code>s[1]</code>
Out:		Out:		Out:		Out:		Out:
a	3.0	7.0		a	3.0	d	4.0	7.0
b	7.0			b	7.0	d	0.3	
c	4.0			c	4.0			
d	4.0							
d	0.3							

- The returned object is either a value, or a subset of the initial series `s`
- Select some data with integer index OR index label
 - **Warning: Work only if the index type is not numeric**

Selection of data

- Filter on DataFrame

In: df
Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df[:2]
Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5

In: df[df['Paris']>1]
Out:

	Paris	Berlin	Madrid
a	3	7	5
c	6	4	8

df.Berlin[df['Berlin']>1]=0
In: df
Out:

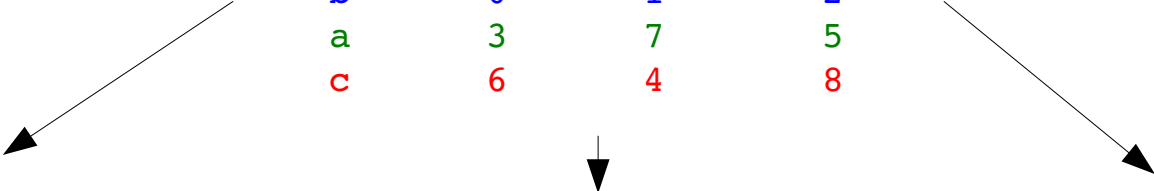
	Paris	Berlin	Madrid
b	0	1	2
a	3	0	5
c	6	0	8

- Output Object: subset of the initial DataFrame

Selection of data

- The indexing field ***ix*** enables to select a subset of the rows and columns from a DataFrame.

```
In: df
Out:
      Paris  Berlin  Madrid
b         0       1       2
a         3       7       5
c         6       4       8
```



```
In: df.ix['a', 'Berlin']
Out:
7
```

```
In: df.ix[['b', 'c'], 'Berlin']
Out:
b     1
c     4
Name: Berlin
```

```
In: df.ix[:, 'Berlin']
Out:
b     1
a     7
c     4
Name: Berlin
```

- Output Object: a value OR a Series subset of the DataFrame

Exercise

Select the rows where 'Rank' = 0

Dropping entries from an axis

- On series or DataFrame, drop a row by his index

```
In: s
Out:
a    3.0
b    7.0
c    4.0
d    4.0
d    0.3
```

```
In: s.drop('d')
Out:
a    3.0
b    7.0
c    4.0
```

```
In: s.drop_duplicates()
Out:
a    3.0
b    7.0
c    4.0
d    0.3
```

- In DataFrame, (default) 'axis=0' refers to (row) index and axis=1 to columns

```
In: df
Out:
   Paris  Berlin  Madrid
b      0       1       2
a      3       7       5
c      6       4       8
```

```
In: df.drop('c')
Out:
   Paris  Berlin  Madrid
b      0       1       2
a      3       7       5
```

```
In: df.drop('Berlin', axis=1)
Out:
   Paris  Madrid
b      0       2
a      3       5
c      6       8
```

Exercise

Drop rows containing 'Rank' = 0

Data alignment


- Series join and align axis to do operations

In: s
Out:
a 3.0
b 7.0
c 4.0
d 4.0
e 0.3

In: s2
Out:
a 0

c 1

f 2



In: s+s2
Out:
a 3
b NaN
c 5
d NaN
e NaN
f NaN

In: s.add(s2, fill_value=0)
Out:
a 3.0
b 7.0
c 5.0
d 4.0
e 0.3
f 2.0

Data alignment

- DataFrame join and align on both axes

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df2

Out:

	Paris	Lisbonne	Madrid
b	0	1	2
e	3	4	5
c	6	7	8
a	9	10	11

In: df+df2

Out:

	Berlin	Lisbonne	Madrid	Paris
a	NaN	NaN	16	12
b	NaN	NaN	4	0
c	NaN	NaN	16	12
e	NaN	NaN	NaN	NaN

In: df.add(df2, fill_value=0)

Out:


	Berlin	Lisbonne	Madrid	Paris
a	7	10	16	12
b	1	1	4	0
c	4	7	16	12
e	NaN	4	5	3

Exercise

- Compute the total amount between the two DataFrame information ('Technical budget' and 'Amount')

Merge, join, concatenate

- Many to one:



In: df1


Out:

	data1	keyLeft
0	0	b
1	1	b
2	2	a
3	3	c
4	4	a
5	5	a
6	6	b

In: df2

Out:

	data2	key
0	0	a
1	1	b
2	2	d



In: `pd.merge(df1,df2, left_on = 'keyLeft', right_on='key', how = 'inner')`

Out:

	data1	keyLeft	data2	key
0	0	b	1	b
1	1	b	1	b
2	6	b	1	b
3	2	a	0	a
4	4	a	0	a
5	5	a	0	a

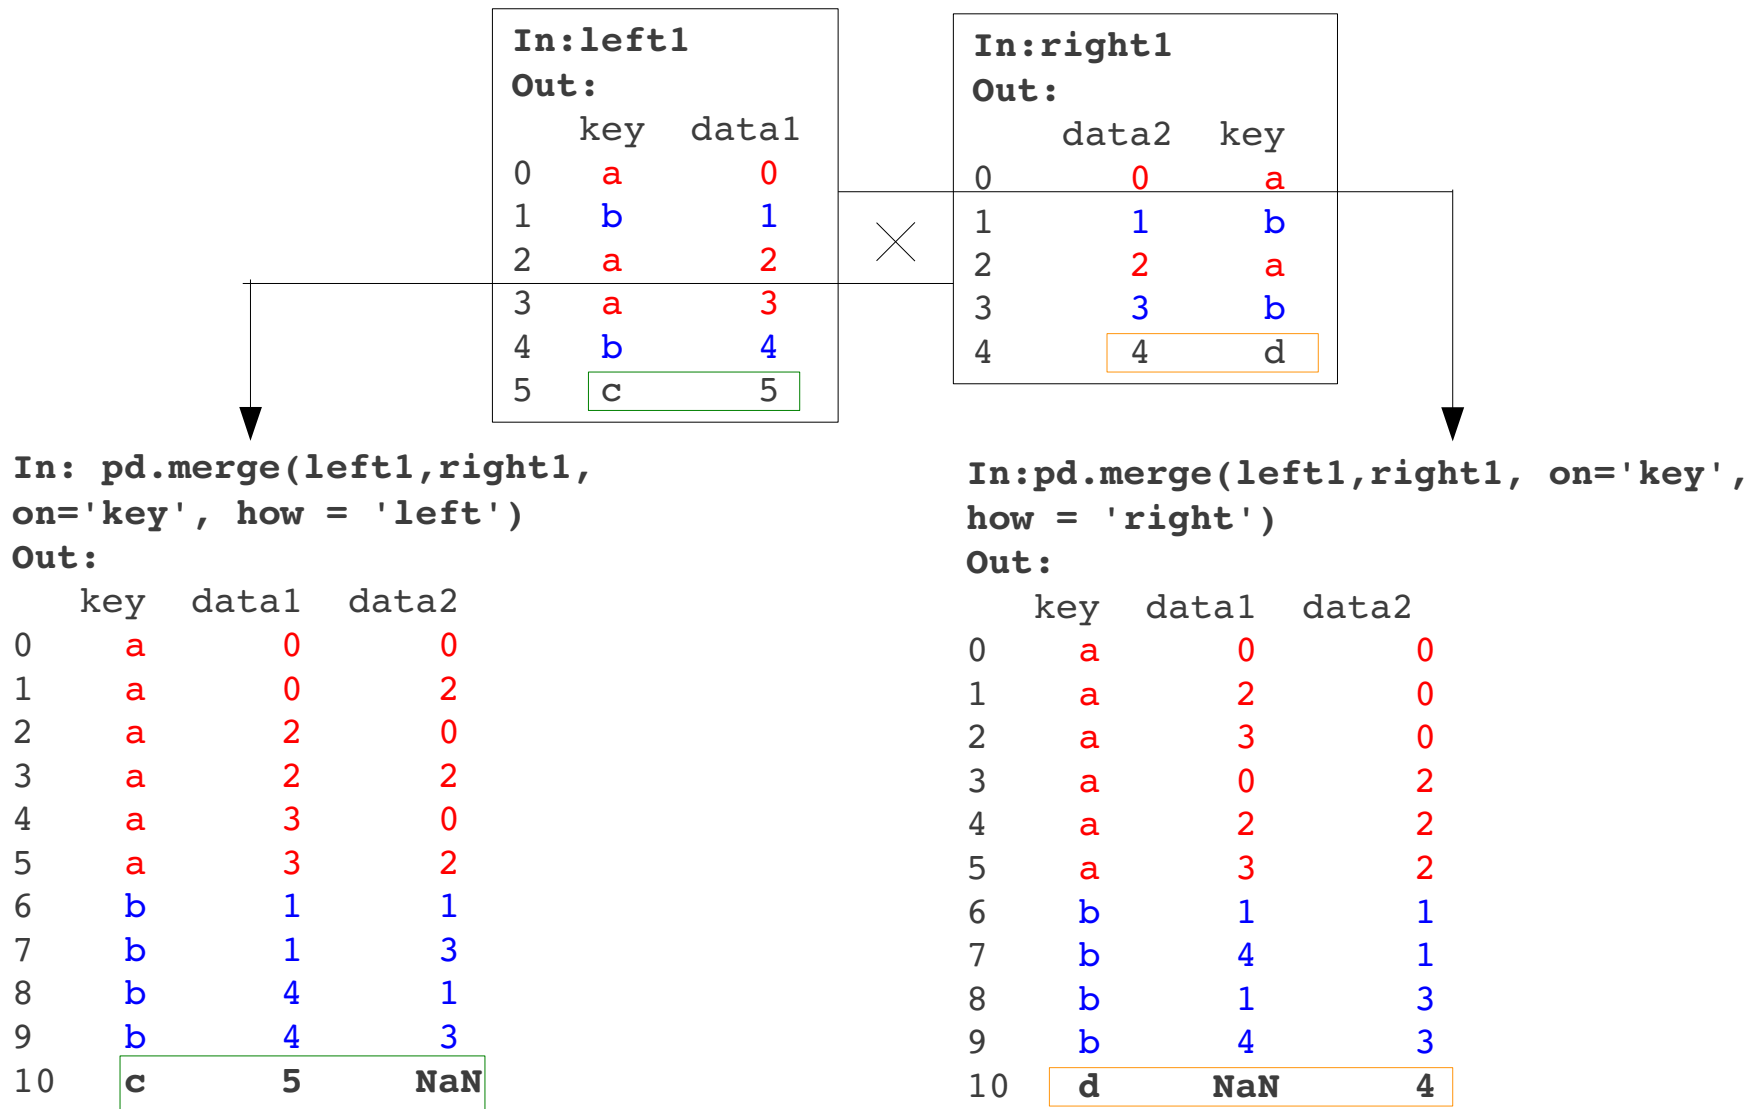
In: `pd.merge(df1,df2, left_on = 'keyLeft', right_on='key', how = 'outer')`

Out:

	data1	keyLeft	data2	key
0	0	b	1	b
1	1	b	1	b
2	6	b	1	b
3	2	a	0	a
4	4	a	0	a
5	5	a	0	a
6	3	c	NaN	NaN
7	NaN	NaN	2	d

Merge, join, concatenate

- Many to many: cartesian product of the rows given a common key



Merge, join, concatenate

Exercise

- Merge the two CSV among the keys [Id, Project] :
- Make the joint considering the intersection

Ranking

- Rank methods on Series and DataFrame among several methods

In: s

Out:

a	3.0
b	7.0
c	4.0
d	4.0
d	0.3

In: s.rank([ascending = True])

Out:

a	2.0
b	5.0
c	3.5
d	3.5
d	1.0

In: s.rank(method='first')

Out:

a	2
b	5
c	3
d	4
d	1

In: s.rank(method='max', ascending=False)

Out:

a	4
b	1
c	3
d	3
d	5

Ranking

- Rank methods on Series and DataFrame

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df.rank()

Out:

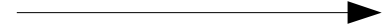
	Paris	Berlin	Madrid
b	1	1	1
a	2	3	2
c	3	2	3



In: df.rank(axis=1)

Out:

	Paris	Berlin	Madrid
b	1	2	3
a	1	3	2
c	2	1	3



- Value = rank in the specified axis

Series ordering/sorting

- Order method: only on Series

In: s


Out:

a	3.0
b	7.0
c	4.0
d	4.0
d	0.3

In: s.order([ascending=True])

Out:

d	0.3
a	3.0
c	4.0
d	4.0
b	7.0

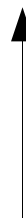


- Sort method by index

In:

s.sort_index(ascending=False)

Out:



d	0.3
d	4.0
c	4.0
b	7.0
a	3.0

DataFrame ordering/sorting

- No order method for DataFrame: specify the axis

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df.sort_index([ascending=True])

Out:

	Paris	Berlin	Madrid
a	3	7	5
b	0	1	2
c	6	4	8

In: df.sort_index(by = 'Berlin')

Out:

	Paris	Berlin	Madrid
b	0	1	2
c	6	4	8
a	3	7	5

In: df.sort_index(axis=1)

Out:

	Berlin	Madrid	Paris
a	7	5	3
b	1	2	0
c	4	8	6

Function application

- Basics operations on Series and DataFrame values

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df.max()

Out:

Paris	6
Berlin	7
Madrid	8

In: df + df.max()

Out:

	Paris	Berlin	Madrid
b	6	8	10
a	9	11	13
c	12	14	16

- Warning: operations are applied among 1D arrays --> the output object is a serie

```
f = lambda x: math.sqrt(x)
```

```
In: df.apply(f)
```

```
Out:
```

```
TypeError: ('only length-1 arrays can be converted to Python  
scalars', u'occurred at index Paris')
```

Function application

- Apply mathematical functions directly on values

```
In: df
```

```
Out :
```

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

```
f = lambda x: math.sqrt(x)
```

```
In: df.applymap(f)
```

```
Out:
```

	Paris	Berlin	Madrid
b	0.000000	1.000000	1.414214
a	1.732051	2.645751	2.236068
c	2.449490	2.000000	2.828427

```
df.Berlin = df['Berlin'].map(f)
```

```
In: df
```

```
Out:
```

	Paris	Berlin	Madrid
b	0	1.000000	2
a	3	2.645751	5
c	6	2.000000	8

Exercise

Assign in a new column 'Total' the sum of the others columns amount values applied with the function $f(x) = x + 0.2 \cdot x$ and sort the table by total value

Computing Descriptive Statistics

- Objects are equipped with a set of common statistical methods.

```
In: df
```

```
Out:
```

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

```
In: df.sum(axis=1)
```

```
Out:
```

b	3
a	15
c	18

```
In: df.describe()
```

```
Out:
```

	Paris	Berlin	Madrid
count	3.0	3.0	3.0
mean	3.0	4.0	5.0
std	3.0	3.0	3.0
min	0.0	1.0	2.0
25%	1.5	2.5	3.5
50%	3.0	4.0	5.0
75%	4.5	5.5	6.5
max	6.0	7.0	8.0

- Covariance and correlation

```
In: df.cov()
```

```
Out:
```

	Paris	Berlin	Madrid
Paris	9.0	4.5	9.0
Berlin	4.5	9.0	4.5
Madrid	9.0	4.5	9.0

```
In: df.corr()
```

```
Out:
```

	Paris	Berlin	Madrid
Paris	1.0	0.5	1.0
Berlin	0.5	1.0	0.5
Madrid	1.0	0.5	1.0

Working on index

- Reindex Series and DataFrame

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

In: df.reindex(['c', 'b', 'a', 'g'])

Out:

	Paris	Berlin	Madrid
c	6	4	8
b	0	1	2
a	3	7	5
g	NaN	NaN	NaN

In: df.reindex(['c', 'b', 'a', 'g'],
fill_value = 14)

Out:

	Paris	Berlin	Madrid
c	6	4	8
b	0	1	2
a	3	7	5
g	14	14	14

In: df.reindex(columns = ['Varsovie', 'Paris', 'Madrid'])

Out:

	Varsovie	Paris	Madrid
b	NaN	0	2
a	NaN	3	5
c	NaN	6	8

Warning: be aware if no duplicate index: `df.index.is_unique`

Hierarchical indexing

- Indices are n-dimensional tables ($n > 1$)
- Easy to build complex datasets

	index		value
b	Paris	▶	0
	Berlin	▶	1
	Madrid	▶	2
a	Paris	▶	3
	Berlin	▶	7
	Madrid	▶	5
c	Paris	▶	6
	Berlin	▶	4
	Madrid	▶	8

In: `df.index`

Out: `MultiIndex`

```
[(b, Paris), (b, Berlin), (b, Madrid),  
 (a, Paris), (a, Berlin), (a, Madrid),  
 (c, Paris), (c, Berlin), (c, Madrid)]
```

- Index are `MultiIndex` objects

Hierarchical indexing

- Build a hierarchical index from DataFrame columns

In: df

Out:

	Paris	Berlin	Madrid
b	0	1	2
a	3	7	5
c	6	4	8

df2 = df.set_index(['Berlin', 'Madrid'])

In: df2

Out:

	Berlin	Madrid	Paris
1	2	0	
7	5	3	
4	8	6	

Diagram illustrating the mapping of values from the original DataFrame to the new hierarchical index structure. The value 5 (originally at index 'a', column 'Madrid') is highlighted in a green box, and the value 3 (originally at index 'a', column 'Paris') is also highlighted in a green box. The value 4 (originally at index 'c', column 'Berlin') is highlighted in an orange box, and the value 6 (originally at index 'c', column 'Paris') is also highlighted in an orange box. Arrows indicate the mapping: a green arrow points from the value 5 to the value 3, and an orange arrow points from the value 4 to the value 6.

- The field **xs** enables to select values from any index level

In: df2.xs(7, level = 0)

Out:

	Paris
Madrid	5

In: df2.xs(8, level = 1)

Out:

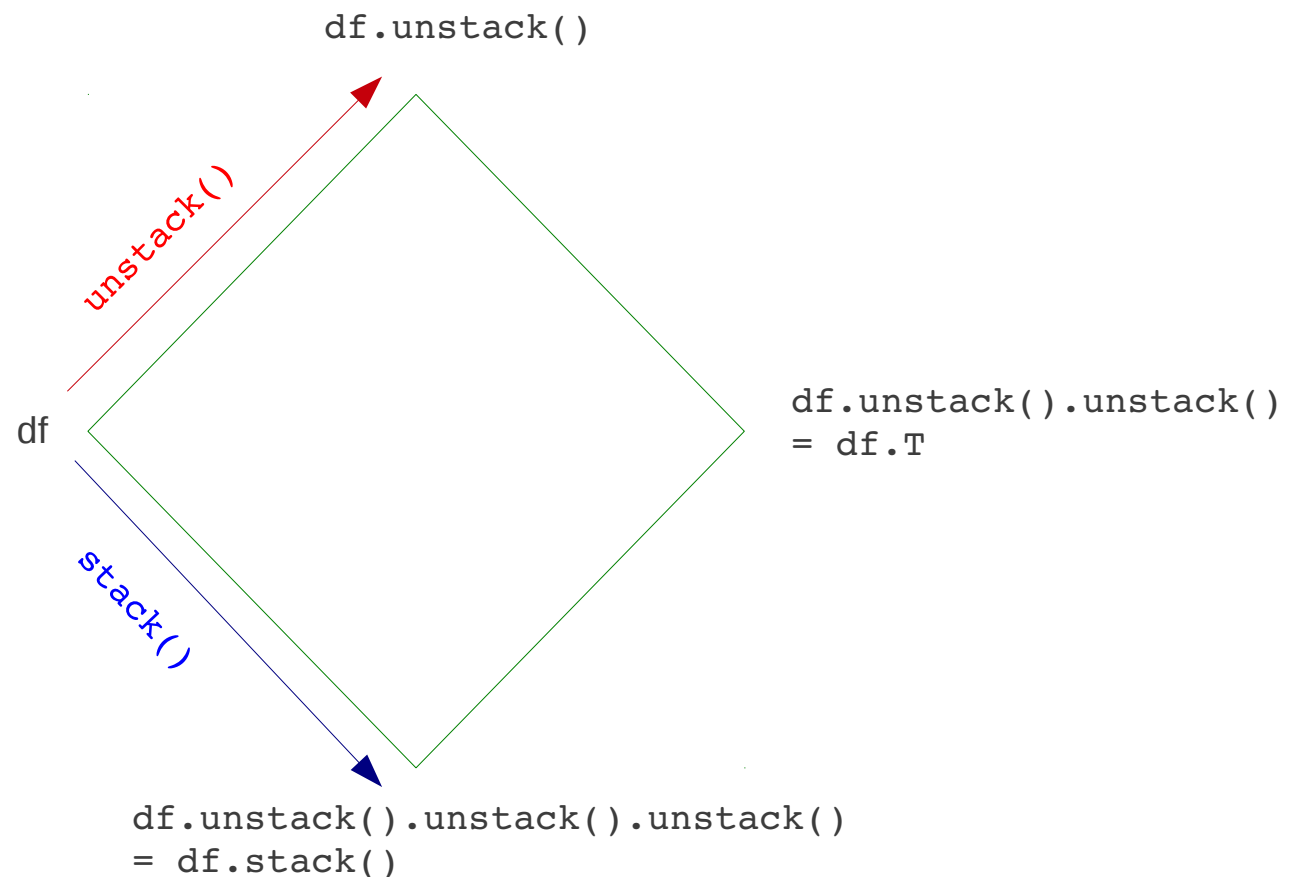
	Paris
Berlin	4

Exercise

Transform the DataFrame dfTot with a hierarchical index: ['Country', 'Id']

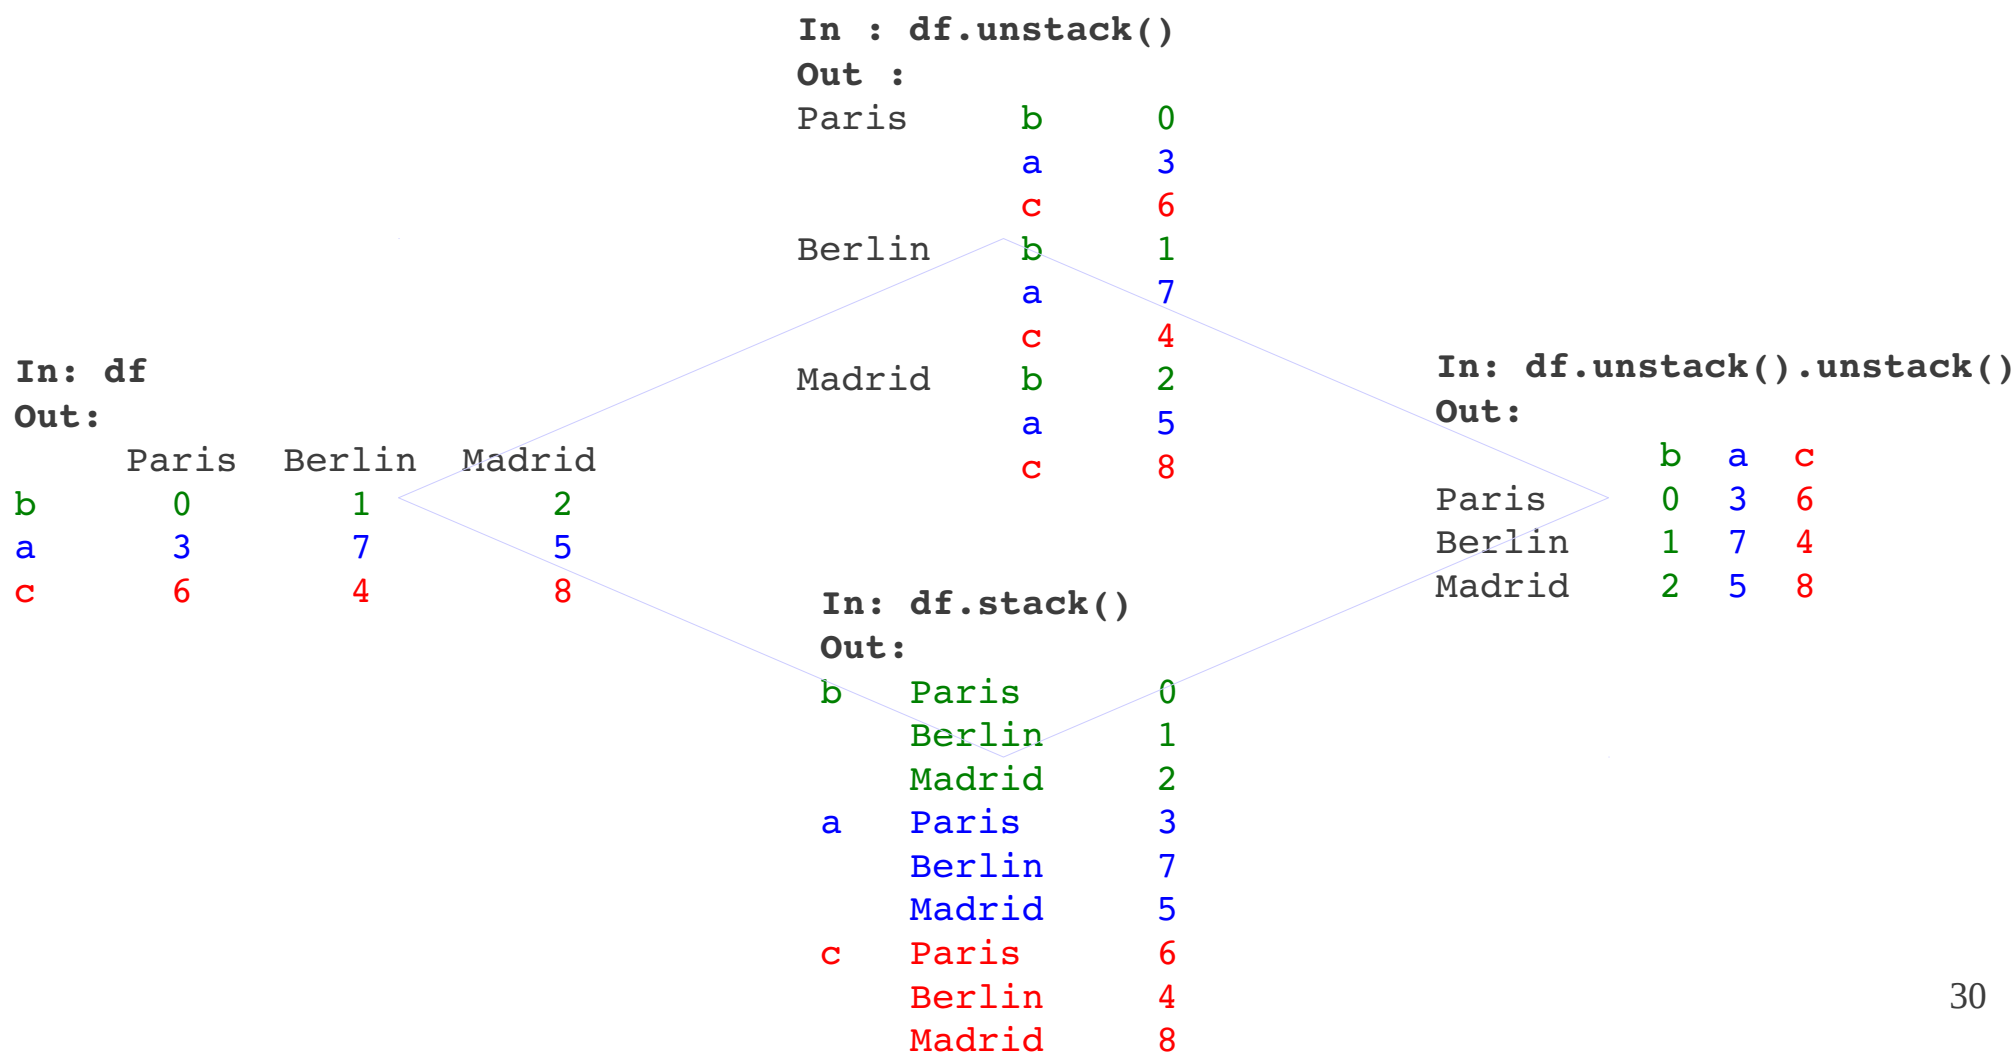
Hierarchical indexing

- Conversion in Series/DataFrame with methods `stack()`/`unstack()`



Hierarchical indexing

- Conversion in Series/DataFrame with methods `stack()`/`unstack()`



The *groupby* Object

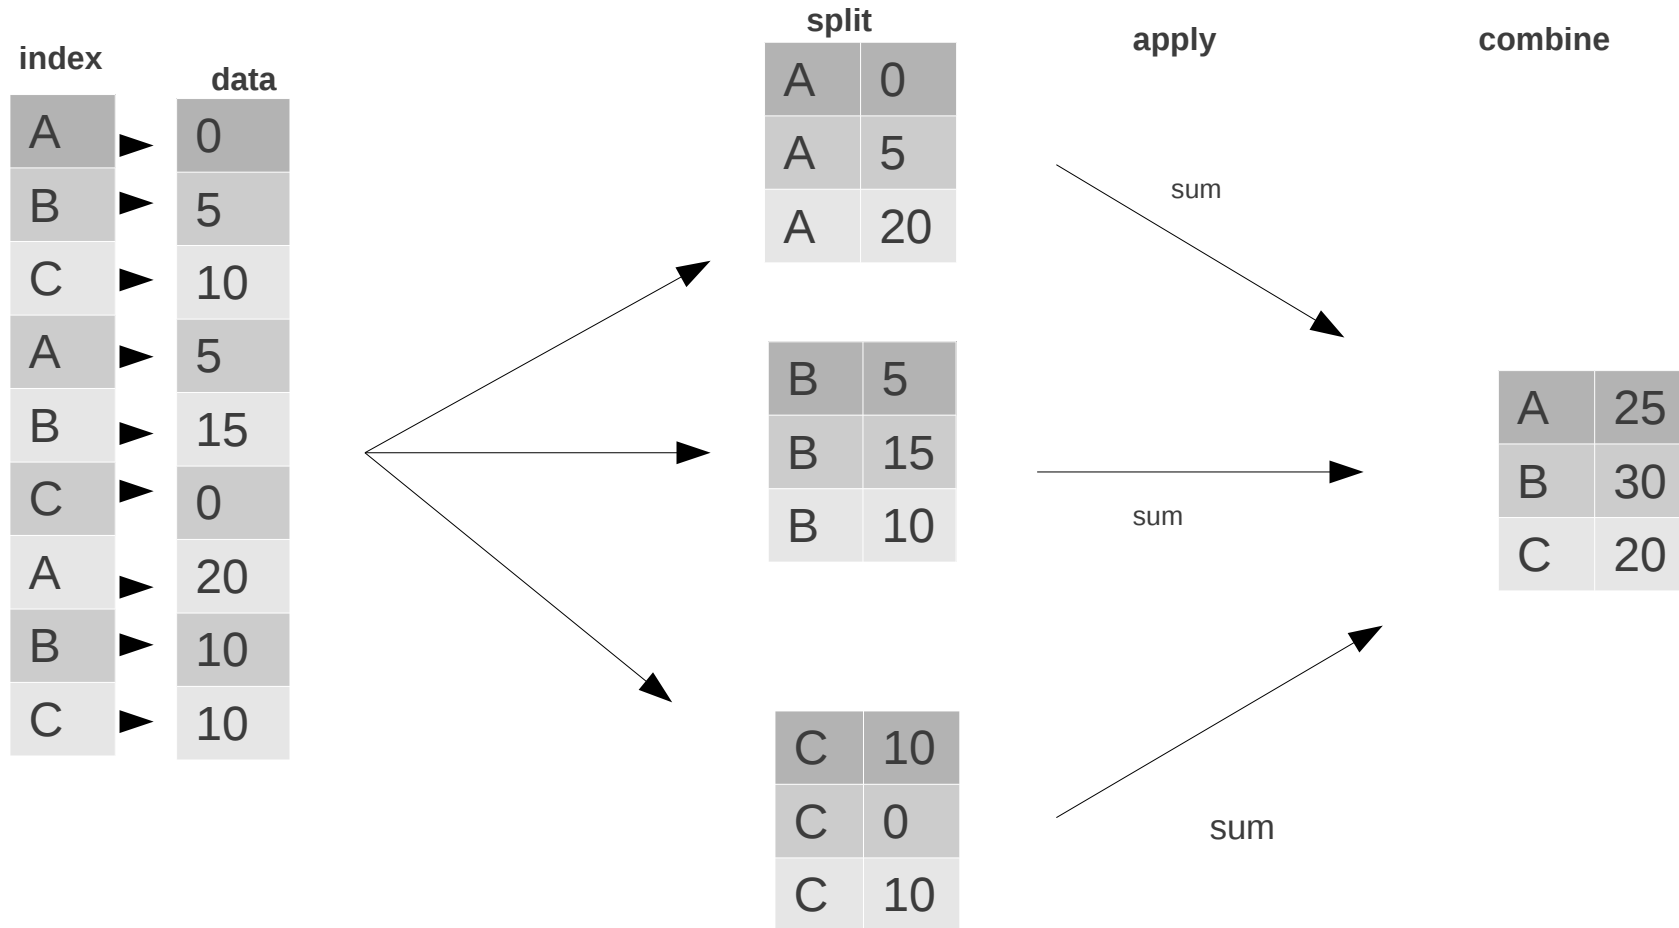


Illustration of a *groupby* process

groupby: a concrete example

In: dfG

Out:

	data1	data2	key1	key2
0	-0.822677	0.120968	a	one
1	0.199444	0.713446	a	two
2	0.054523	-0.530082	b	one
3	-1.087544	-1.952220	b	two
4	0.591362	-0.446848	a	one

In: group = dfG.groupby(['key1', 'key2'])
group

Out:

<pandas.core.groupby.DataFrameGroupBy
object at 0x3960f90>

- All operations are possible from the groupby

In: group[data1].mean()

Out:

	key1	key2	data1
a	one	-0.115657	
	two	0.199444	
b	one	0.054523	
	two	-1.087544	

In: group.size()

Out:

key1	key2	
a	one	2
	two	1
b	one	1
	two	1

Exercise

Give the mean 'Rank' by 'Id' using the groupby object
Drop the rows which contains duplicated Id

Draw the graph

Exercise

- Our DataFrame is now cleaner and well ordered
 - Draw the associated graph:
 - Nodes: Project item
 - Edges: same country between two nodes
 - Set the DataFrame index with the referenced `tlp.node` object
 - Create a new DataFrame with all the graphic properties (at least `viewLayout`) of each node
 - using this DataFrame, draw the nodes on a line starting at the first node's position

Conclusion

- Manipulate data easily and fastly
- Intuitive representation
- N-hierarchical index and groupby: most powerful tool of pandas
- Statistics methods and calculs based on R language
- Nearly impossible to combine or adjust misaligns data

